

## Hardware Drivers

### Contents

1. [ADC](#)
2. [Clock](#)
  - 2.1. [Driver configuration](#)
  - 2.2. [Using the clock driver](#)
    - 2.2.1. [Controlling the HFCLK](#)
    - 2.2.2. [Controlling the LFCLK \(without a SoftDevice only\)](#)
    - 2.2.3. [Controlling the calibration of the LFCLK \(without a SoftDevice and using RC as source clock\)](#)
3. [GPIOTE](#)
  - 3.1. [TASK/EVENT channels allocation](#)
  - 3.2. [Driver configuration](#)
  - 3.3. [Driver initialization](#)
  - 3.4. [Controlling output pins](#)
    - 3.4.1. [Initialization](#)
    - 3.4.2. [Manually controlled output pins](#)
    - 3.4.3. [Task-controlled output pins](#)
  - 3.5. [Controlling input pins](#)
    - 3.5.1. [Initialization](#)
    - 3.5.2. [Usage](#)
    - 3.5.3. [Working with a SoftDevice](#)
4. [I2S](#)
  - 4.1. [Driver configuration](#)
  - 4.2. [Using the I2S driver](#)
5. [LPCOMP](#)
6. [PDM](#)
  - 6.1. [Driver configuration](#)
  - 6.2. [Using the PDM driver](#)
7. [PPI](#)
8. [PWM](#)
  - 8.1. [Driver configuration](#)
  - 8.2. [Using the PWM driver](#)
9. [QDEC](#)
10. [RNG](#)
11. [SAADC](#)
  - 11.1. [Driver configuration](#)
  - 11.2. [Using the SAADC driver](#)
    - 11.2.1. [Blocking mode](#)
    - 11.2.2. [Non-blocking mode](#)
    - 11.2.3. [Limits](#)
12. [SPI master](#)
  - 12.1. [Driver configuration](#)
  - 12.2. [Basic usage](#)
  - 12.3. [Advanced usage](#)

- [12.3.1. Starting a transfer from PPI](#)
    - [12.3.2. Repeated transfers](#)
  - [12.4. Events](#)
- [13. SPI slave](#)
  - [13.1. Driver configuration](#)
  - [13.2. Using the SPI slave driver](#)
- [14. SWI](#)
  - [14.1. Usage](#)
    - [14.1.1. Initialization](#)
    - [14.1.2. Allocation](#)
    - [14.1.3. Event handler and flags](#)
    - [14.1.4. Special options](#)
  - [14.2. Example](#)
  - [14.3. Usage with a SoftDevice](#)
- [15. Timer](#)
- [16. TWI master](#)
  - [16.1. Initialization](#)
  - [16.2. Basic usage](#)
  - [16.3. Advanced usage](#)
    - [16.3.1. Starting a transfer from PPI](#)
    - [16.3.2. Repeated transfers](#)
  - [16.4. Events](#)
- [17. TWI slave](#)
  - [17.1. Driver configuration](#)
  - [17.2. Modes of operation](#)
    - [17.2.1. Synchronous mode](#)
    - [17.2.2. Asynchronous mode](#)
- [18. UART](#)
  - [18.1. Driver configuration](#)
  - [18.2. Using the UART driver](#)
    - [18.2.1. Blocking mode](#)
    - [18.2.2. Non-blocking mode](#)
    - [18.2.3. Enabling RX without providing data buffer](#)
- [19. WDT](#)

nRF5 SDK v11.0.0-2.alpha

## Hardware Drivers

[ADC](#)

nRF51 only

[Clock](#)

[GPIOTE](#)

[I2S](#)

nRF52 only

[LPCOMP](#)[PDM](#) *nRF52 only*[PPI](#)[PWM](#) *nRF52 only*[QDEC](#)[RNG](#)[SAADC](#) *nRF52 only*[SPI master](#)[SPI slave](#)[SWI](#)[Timer](#)[TWI master](#)[TWI slave](#) *nRF52 only*[UART](#)[WDT](#)

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

*1. nRF5 SDK v11.0.0-2.alpha*

## ADC

*This information applies to the **nRF51 Series** only.*

The ADC hardware access layer (HAL) provides basic APIs for accessing the registers of the analog-to-digital converter (ADC). See the API documentation for the [ADC HAL](#) for details.

The [ADC HAL Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

2. nRF5 SDK v11.0.0-2.alpha

## Clock

The clock driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the clock. See the API documentation for the [Clock HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [Clock driver](#) for details.

The clock driver is responsible for starting the high-accuracy, high-frequency clock source and the low-frequency clock (without a SoftDevice). It also performs calibration without a SoftDevice or internal RC as an LFCLK source.

### Driver configuration

Enable the driver in the configuration file `nrf_drv_config.h` before using it in an application. The clock driver supports only the static, compile-time configuration due to the possibility of other modules attempting to initialize it. The configuration includes the high-frequency clock frequency, the low-frequency clock source, and the interrupt priority.

```
#define CLOCK_ENABLED 1
```

Enabling the clock will initialize the [Clock driver](#) to its default state:

- High-accuracy, high-frequency clock HFCLK is off
- Low-frequency clock LFCLOCK is off (without a SoftDevice)
- Calibration has not started (without a SoftDevice only)

### Using the clock driver

The [Clock Example](#) provides sample code that you can use to quickly get started.

### Controlling the HFCLK

If a module needs a high-precision clock, it can request the clock driver to start it. If a callback is provided, it will be called after the request is handled. If the requested clock was already started when [nrf\\_drv\\_clock\\_hfclk\\_request](#) is called, the callback will be called from the context of this function. Otherwise, it will be called from the context of the clock interrupt handler. To check if the HFCLK is on, call [nrf\\_drv\\_clock\\_hfclk\\_is\\_running](#).

When the module is finished with the HFCLK, you should release it by calling [nrf\\_drv\\_clock\\_hfclk\\_release](#). The clock is turned off if no other module is requesting it.

The SoftDevice API is used in the driver if a SoftDevice is present.

## Controlling the LFCLK (without a SoftDevice only)

If a module needs a low-frequency clock, it can request the clock driver to start it. If a callback is provided, it will be called after the request is handled. If the requested clock was already started when [nrf\\_drv\\_clock\\_lfclk\\_request](#) is called, the callback will be called from the context of this function. Otherwise, it will be called from the context of the clock interrupt handler. To check if the LFCLK is on, call [nrf\\_drv\\_clock\\_lfclk\\_is\\_running](#).

When the module is finished with the LFCLK, you should release it by calling [nrf\\_drv\\_clock\\_lfclk\\_release](#). The clock is turned off if no other module is requesting it.

The SoftDevice API is used in the driver if a SoftDevice is present.

## Controlling the calibration of the LFCLK (without a SoftDevice and using RC as source clock)

The internal RC can be used as source for the LFCLK, but it is not as precise as a crystal. It must be periodically calibrated against a high-frequency crystal. The clock module supports autonomous calibration. It also has a built-in low-power timer dedicated to calibration.

The calibration process consists of three phases:

- Delay (using a clock calibration timer); delay can be set to 0 to start calibration immediately
- Request the high-accuracy HFCLK
- Perform calibration

Call [nrf\\_drv\\_clock\\_calibration\\_start](#) to start the calibration process. Parameters specify the delay (given in 0.25 ms units) and an optional callback to be called once calibration is completed.

To check if there is an ongoing calibration, call [nrf\\_drv\\_clock\\_is\\_calibrating](#).

At any time, [nrf\\_drv\\_clock\\_calibration\\_abort](#) can be called to stop calibration. If a callback was provided for [nrf\\_drv\\_clock\\_calibration\\_start](#), the callback is executed when the process stops, with the event type set to [NRF\\_DRV\\_CLOCK\\_EVT\\_CAL\\_ABORTED](#).

Periodic calibration can be achieved by calling [nrf\\_drv\\_clock\\_calibration\\_start](#) from the calibration completion handler:

```
void clock_handler(nrf_drv_clock_evt_type_t event)
{
    if (event == NRF_DRV_CLOCK_EVT_CAL_DONE)
    { ret = nrf_drv_clock_calibration_start(10, clock_handler); }
}
```

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

3. nRF5 SDK v11.0.0-2.alpha

## GPIOTE

The GPIOTE driver controls the GPIO and GPIOTE peripherals and configures and controls output and input pins. Output pins can be configured to be controlled manually or by a GPIOTE task. Input pins can be configured to be controlled in different modes:

- High accuracy: An independent GPIOTE event is used to detect changes on the pin. If a pin is configured to be controlled in this mode, a high frequency clock is required.
- Low accuracy/low power: A PORT event senses level changes on the pin. One PORT event can be used for multiple pins. Therefore, it is not that accurate and cannot be used to track high-speed pin changes. On the other hand, if a pin is configured in this mode, the system can turn off the high frequency clock when sleeping.

### TASK/EVENT channels allocation

The number of TASK/EVENT channels that can be used to drive an output pin using a GPIOTE task or to generate an event on input pin change is limited. The driver manages those channels. The user cannot control which channels are used.

When [nrf\\_drv\\_gpiote\\_in\\_init](#) or [nrf\\_drv\\_gpiote\\_out\\_init](#) is called, the driver allocates a channel from the pool. If all channels are already allocated, the functions return an error code. Note that [nrf\\_drv\\_gpiote\\_out\\_init](#) allocates a channel only if [nrf\\_drv\\_gpiote\\_out\\_config\\_t](#) indicates that the pin will be controlled by a task.

When [nrf\\_drv\\_gpiote\\_in\\_uninit](#) or [nrf\\_drv\\_gpiote\\_out\\_uninit](#) is called, the driver frees the allocated channel.

### Driver configuration

To ensure that the driver configuration is known to all modules, only static configuration is supported. Therefore, [nrf\\_drv\\_gpiote\\_init](#) has no input parameters. The static configuration parameters are located in the [nrf\\_drv\\_config.h](#) file.

The configuration parameters include:

- Interrupt priority level
- Maximum amount of low-power events (value needed to statically allocate array for user handlers)

### Driver initialization

The GPIOTE driver is a shared resource that can be used by multiple modules in an application. Therefore, it can be initialized only once. If a module is using the driver, it must check if it has already been initialized by calling the function [nrf\\_drv\\_gpiote\\_is\\_init](#). If this function returns false, the module must initialize the driver by calling the function [nrf\\_drv\\_gpiote\\_init](#).

The following code example shows how to initialize the driver:

```
uint32_t err_code;  
if(!nrf_drv_gpiote_is_init())  
{  
    err_code = nrf_drv_gpiote_init();  
}
```

## Controlling output pins

After initialization, output pins can be configured to be controlled manually or by a GPIOTE task.

### Initialization

Output pins are initialized by calling the function [nrf\\_drv\\_gpiote\\_out\\_init](#). The argument of type [nrf\\_drv\\_gpiote\\_out\\_config\\_t](#) specifies the configuration of the output pin. You can use the following macros to set the configuration:

- [GPIOTE\\_CONFIG\\_OUT\\_SIMPLE](#): Indicates that the pin will be driven by driver function calls and not by a task. As parameter, you can specify the initial pin state.
- [GPIOTE\\_CONFIG\\_OUT\\_TASK\\_LOW](#): Indicates that a task will be used to clear the pin. The pin is initially set.
- [GPIOTE\\_CONFIG\\_OUT\\_TASK\\_HIGH](#): Indicates that a task will be used to set the pin. The pin is initially cleared.
- [GPIOTE\\_CONFIG\\_OUT\\_TASK\\_TOGGLE](#): Indicates that a task will be used to toggle the pin. The pin is initially set. As parameter, you can specify the initial pin state.

The configuration can be modified before calling [nrf\\_drv\\_gpiote\\_out\\_init](#).

The following code example configures the 10th pin to be used by the task that sets the pin. The pin is initially set.

```
nrf_drv_gpiote_out_config_t config = GPIOTE_CONFIG_OUT_TASK_HIGH;  
config.init_set = true;  
err_code = nrf_drv_gpiote_out_config(10, &config);
```

## Manually controlled output pins

When a pin is configured as non-task pin, it can be controlled by calling the following functions:

- [nrf\\_drv\\_gpiote\\_out\\_set](#)
- [nrf\\_drv\\_gpiote\\_out\\_clear](#)
- [nrf\\_drv\\_gpiote\\_out\\_toggle](#)

The driver asserts if the functions are called on a pin that is configured as task pin.

## Task-controlled output pins

When a pin is configured as task pin, it can be controlled using the CPU by writing 1 to the relevant TASK register, or using PPI. Before a pin can be controlled, it must be enabled by calling the function [nrf\\_drv\\_gpiote\\_out\\_task\\_enable](#).

When a pin is initialized as task pin, it allocates one of the available TASK/EVENT channels. Therefore, you must call [nrf\\_drv\\_gpiote\\_out\\_uninit](#) when the pin is no longer used.

The following code example shows how to use a task pin to be toggled using a task:

```
nrf_drv_gpiote_out_config_t config = GPIOTE_CONFIG_OUT_TASK_TOGGLE
    (false);
err_code = nrf_drv_gpiote_out_init(10, &config);
if (err_code != NRF_SUCCESS)
{
    // handle error condition
}
ppi_task_addr = nrf_drv_gpiote_out_task_addr_get(10);
// Configure PPI using ppi_task_address or use the address to manually
    control the task.

nrf_drv_gpiote_out_task_enable(10);

// Task is enabled and can be used by PPI.

nrf_drv_gpiote_out_task_disable(10);

// Task is disabled and cannot be controlled.

nrf_drv_gpiote_out_task_enable(10);

// Task is enabled again and can be controlled using ppi_task_addr.

nrf_drv_gpiote_out_uninit(10);

// Pin is uninitialized. Task is disabled and TASK/EVENT channel is
    freed.
```

## Controlling input pins

The GPIOTE peripheral can generate events on pin level change using the following methods:

- Using a dedicated IN\_EVENT: When using a dedicated IN\_EVENT, one pin is assigned to an event. Only a limited number of pins can be sensed this way because there is a limited number of TASK/EVENT channels. When an IN\_EVENT is enabled, it requests the high frequency clock. Therefore, this method is not suitable for scenarios where the system sleeps for a long period of time, expecting to be waken up by a pin level change. This method supports detection of toggling, low-to-high, and high-to-low transition.
- Using a PORT event: A PORT event is a single event that is triggered whenever the value of any of the input pins that have sensing enabled changes. Pins can be configured to sense low-to-high or high-to-low transition. Toggle sensing is emulated in the driver by changing the sense configuration whenever a PORT event for a given pin is detected. Using a PORT event requires only a low frequency clock. Therefore, this method is



suitable for scenarios where the application expects to be inactive for a long period of time, waiting for a pin level change. The event can be shared by multiple pins; therefore it should be used for sensing slowly changing signals only.

## Initialization

Input pins are initialized by calling the function [nrf\\_drv\\_gpiote\\_in\\_init](#). The argument of type [nrf\\_drv\\_gpiote\\_in\\_config\\_t](#) specifies the configuration of the input pin. You can use the following macros to set the configuration:

- [GPIOTE\\_CONFIG\\_IN\\_SENSE\\_LOTOHI](#): Indicates that the input pin will generate an event when the input value is changed to high. As parameter, you can specify if the input should have high accuracy or not.
- [GPIOTE\\_CONFIG\\_IN\\_SENSE\\_HITOLO](#): Indicates that the input pin will generate an event when the input value is changed to low. As parameter, you can specify if the input should have high accuracy or not.
- [GPIOTE\\_CONFIG\\_IN\\_SENSE\\_TOGGLE](#): Indicates that the input pin will generate an event when the input value is changed. As parameter, you can specify if the input should have high accuracy or not.

The configuration can be modified before calling [nrf\\_drv\\_gpiote\\_in\\_init](#). Note that macros are not setting any pull.

During the initialization of an input pin, the user must provide an event handler. The event handler will be called by the driver when the configured transition occurs and the interrupt is enabled.

The following code example configures the 10th pin to have pull-up and generate an event on pin toggling:

```
nrf_drv_gpiote_in_config_t config = GPIOTE_CONFIG_IN_SENSE_TOGGLE;  
config.pull = NRF_GPIO_PIN_PULLUP;  
err_code = nrf_drv_gpiote_in_config(10, &config, pin_event_handler);
```

## Usage

From a software point of view, the input sensing functionality is independent of the pin mode. If a high-accuracy pin is used, it can detect signals that are changed faster. Turning off high accuracy, on the other hand, enables the low-power feature.

Events that are generated by an input pin level change can be used to trigger events to be utilized by PPI, and optionally they can trigger a user event handler.

To start sensing an input pin, it must be enabled.

Pin input values can always be checked by calling the function [nrf\\_drv\\_gpiote\\_in\\_is\\_set](#).

The following code example shows how to use a pin to detect toggling using high accuracy:

```

nrf_drv_gpiote_in_config_t config = GPIOTE_CONFIG_IN_SENSE_TOGGLE
    (true);
err_code = nrf_drv_gpiote_in_init(10, &config, pin_event_handler);
if (err_code != NRF_SUCCESS)
{
    // handle error condition
}
ppi_event_addr = nrf_drv_gpiote_in_event_addr_get(10);
// Configure PPI using ppi_event_address.

// Enable event, interrupt is not enabled.
nrf_drv_gpiote_in_event_enable(10, false);

// Event is enabled and can be used by PPI.

nrf_drv_gpiote_in_event_disable(10);

// Task is disabled and cannot be controlled.

// Event and corresponding interrupt are enabled.
nrf_drv_gpiote_in_event_enable(10, true);

// Event is enabled again and can be used by PPI.
// Additionally, any toggling sensed will trigger interrupt

nrf_drv_gpiote_in_uninit(10);

// Pin is uninitialized. Event is disabled and TASK/EVENT channel is
    freed.

```

## Working with a SoftDevice

When a SoftDevice is enabled, it can interrupt the application at any time for a certain amount of time. This can lead to the situation where some pin level changes are missed. If the application must track an input pin that can change its level frequently, PPI should be used with a high-accuracy event together with a TIMER in counter mode to count the detected transitions.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

4. nRF5 SDK v11.0.0-2.alpha

## I2S

*This information applies to the nRF52 Series only.*

The Inter-IC Sound (I2S) interface driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the I2S peripheral. See the API documentation for the [I2S HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [I2S driver](#) for details.

Key features include:

- Three modes of operation:
  - Reception (RX) only
  - Transmission (TX) only
  - Both directions simultaneously
- Data exchange through a single data handler that is called by the driver when new data is received or new data is needed for transmission

The [I2S Loopback Example](#) provides sample code that you can use to quickly get started.

## Driver configuration

The configurable parameters include:

- Pins to be used for I2S peripheral signals
- Operating mode (master or slave)
- Format of frames, used channels (stereo, left or right)
- Width and alignment of samples
- Clock settings (defining the sampling frequency in master mode)

The default configuration for the I2S driver is located in `nrf_drv_config.h`. This configuration is used when [nrf\\_drv\\_i2s\\_init](#) is called with a NULL pointer to the configuration structure.

The driver must be explicitly enabled in `nrf_drv_config.h` before it can be used. A proper data handling function must be provided during the initialization.

## Using the I2S driver

First, define the data handling function for I2S transfers. For example:

```
static void data_handler(uint32_t const * p_data_received,
                        uint32_t      * p_data_to_send,
                        uint16_t        number_of_words)
{
    if (p_data_received != NULL)
    {
        // Process the received data.
    }

    if (p_data_to_send != NULL)
    {
        // Provide the data to be sent.
    }
}
```

```

    }
}

```

Next, call [nrf\\_drv\\_i2s\\_init](#) to initialize and configure the instance. See the [nrf\\_drv\\_i2s\\_config\\_t](#) structure documentation for possible configuration options.

For example:

```

uint32_t err_code;

nrf_drv_i2s_config_t config = NRF_DRV_I2S_DEFAULT_CONFIG;
config.mck_setup = NRF_I2S_MCK_32MDIV21;
config.ratio      = NRF_I2S_RATIO_96X;
err_code = nrf_drv_i2s_init(&config, data_handler);
if (err_code != NRF_SUCCESS)
{
    // Initialization failed. Take recovery action.
}

```

Start the transfer by calling [nrf\\_drv\\_i2s\\_start](#) and providing a proper buffer for a given direction if it should be enabled. See the following examples:

- RX only:

```
nrf_drv_i2s_start(m_buffer_rx, NULL, I2S_BUFFER_SIZE, 0);
```

- TX only:

```
nrf_drv_i2s_start(NULL, m_buffer_tx, I2S_BUFFER_SIZE, 0);
```

- Both directions:

```
nrf_drv_i2s_start(m_buffer_rx, m_buffer_tx, I2S_BUFFER_SIZE, 0);
```

If both RX and TX are enabled, you can use the [NRF\\_DRV\\_I2S\\_FLAG\\_SYNCHRONIZED\\_MODE](#) flag to force a common call to the data handler for RX and TX data. Normally, the driver calls the data handler as soon as the event for a given direction is generated by the I2S peripheral. In synchronized mode, the call is deferred until events for both directions occur. This mode is useful, for example, when received data should be processed and then transmitted.

At any time, the transfer can be stopped immediately by calling [nrf\\_drv\\_i2s\\_stop](#).

When the I2S driver is no longer needed or its configuration must be changed, call [nrf\\_drv\\_i2s\\_uninit](#).

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

5. nRF5 SDK v11.0.0-2.alpha

## LPCOMP

The Low Power Comparator (LPCOMP) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the LPCOMP. All functions in this layer are implemented as inline functions. See the API documentation for the [LPCOMP HAL](#) for details.

The driver layer provides support for single instances and APIs on a higher level than the HAL. See the API documentation for the [LPCOMP driver](#) for details.

The [LPCOMP Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

6. nRF5 SDK v11.0.0-2.alpha

## PDM

*This information applies to the nRF52 Series only.*

The PDM driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the PDM peripheral. See the API documentation for the [PDM HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [PDM driver](#) for details.

Key features include:

- Asynchronous conversion using DMA.
- Mono or stereo sampling.
- Adjustable gain levels.
- PCM output.

### Driver configuration

The PDM default configuration is located in `nrf_drv_config.h`. The PDM peripheral must be explicitly enabled in `nrf_drv_config.h` before the driver can be used. You must provide a valid configuration structure and an event handler function during initialization. The driver requires two sample buffers that are filled alternately.

The interface is initialized by calling the [nrf\\_drv\\_pdm\\_init](#) function. The driver configuration includes:

- Sampling mode.
- PDM clock frequency.
- Gain levels.
- Two sample buffers.
- Interrupt priority level.

Example:

```
uint16_t buff[2][512]

void pdm_event_handler(uint32_t * p_buffer, uint16_t length)
{
    // Process the data.
}

/* ... */

nrf_drv_pdm_config_t pdm_cfg = NRF_DRV_PDM_DEFAULT_CONFIG(CLK_PIN,
    DOUT_PIN,
    buff[0], buff
    [1], 512);
ret_code_t err_code = nrf_drv_pdm_init(&pdm_cfg, &pdm_event_handler);
```

## Using the PDM driver

After initialization, sampling is started by calling the [nrf\\_drv\\_pdm\\_start](#) function. When the current buffer is filled with data, the event handler function is called with the buffer pointer and size as arguments, and the driver starts filling the second buffer. The output PCM samples are stored as 16-bit signed integers. If stereo mode is selected, the buffer will contain alternating left and right channel samples. Sampling runs continuously until [nrf\\_drv\\_pdm\\_stop](#) is called.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

7. nRF5 SDK v11.0.0-2.alpha

## PPI

The Programmable Peripheral Interconnect (PPI) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the PPI. All functions in this layer are implemented as inline functions. See the API documentation for the [PPI HAL](#) for details.

The driver layer provides support for the PPI module and APIs on a higher level than the HAL. See the API documentation for the [PPI driver](#) for details.

The [PPI Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

8. nRF5 SDK v11.0.0-2.alpha

## PWM

*This information applies to the nRF52 Series only.*

The Pulse Width Modulation (PWM) module driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the PWM peripheral. See the API documentation for the [PWM HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [PWM driver](#) for details.

Key features include:

- Multi-instance support
- Two kinds of playback:
  - Simple: a single sequence of duty cycle values
  - Complex: two concatenated sequences
- Sequences are played back a specified number of times
- Optionally, the whole playback can be repeated in a loop
- A user-defined event handler can be used to perform additional actions when a sequence or the whole playback is finished
- If no event handler is used, playbacks can be carried out without involving the CPU

The [PWM Driver Example](#) provides sample code that you can use to quickly get started.

## Driver configuration

The configurable parameters include:

- Pins to be used as outputs for PWM channels (up to 4 per instance)
- Base clock frequency and top counter value (these parameters jointly determine the PWM period and resolution)
- Mode of loading the sequence values from RAM and distributing them to compare registers:
  - Common: One value is used in all channels.

- Grouped: One value is used in channels 0 and 1, and another one in channels 2 and 3.
- Individual: Separate values are used in each channel.
- Waveform: Similar to Individual mode, but only three channels are used. The fourth value is loaded into the pulse generator counter as its top value.
- Mode of advancing the active sequence (automatic or triggered via a task)

The default configuration for the PWM driver is located in `nrf_drv_config.h`. This configuration is used when [nrf\\_drv\\_pwm\\_init](#) is called with a NULL pointer to the configuration structure.

Before it can be used for a particular instance of the peripheral, the driver must be explicitly enabled for this instance in `nrf_drv_config.h`.

## Using the PWM driver

First, call [nrf\\_drv\\_pwm\\_init](#) to initialize and configure the driver for a given instance of the peripheral. See the [nrf\\_drv\\_pwm\\_config\\_t](#) structure documentation for possible configuration options.

For example:

```
uint32_t err_code;
nrf_drv_pwm_config_t const config0 =
{
    .output_pins =
    {
        BSP_LED_0 | NRF_DRV_PWM_PIN_INVERTED, // channel 0
        NRF_DRV_PWM_PIN_NOT_USED,             // channel 1
        NRF_DRV_PWM_PIN_NOT_USED,             // channel 2
        NRF_DRV_PWM_PIN_NOT_USED,             // channel 3
    },
    .base_clock = NRF_PWM_CLK_1MHz,
    .count_mode = NRF_PWM_MODE_UP,
    .top_value   = 1000,
    .load_mode   = NRF_PWM_LOAD_COMMON,
    .step_mode   = NRF_PWM_STEP_AUTO
};
err_code = nrf_drv_pwm_init(&m_pwm0, &config0, NULL);
if (err_code != NRF_SUCCESS)
{
    // Initialization failed. Take recovery action.
}
```

Next, define a sequence of duty cycle values using the [nrf\\_pwm\\_sequence\\_t](#) structure. For example:

```
static nrf_pwm_values_common_t seq_values[] =
{
```



```

    0, 200, 400, 600, 800, 1000
};
nrf_pwm_sequence_t const seq =
{
    .values.p_common = seq_values,
    .length           = NRF_PWM_VALUES_LENGTH(seq_values),
    .repeats          = 0,
    .end_delay        = 0
};

```

Start the playback of the sequence by calling [nrf\\_drv\\_pwm\\_simple\\_playback](#).

You can play the sequence several times (for example, 3 times):

```
nrf_drv_pwm_simple_playback(&m_pwm0, &seq, 3, 0);
```

Additionally, the playback can be repeated in a loop:

```
nrf_drv_pwm_simple_playback(&m_pwm0, &seq, 3, NRF_DRV_PWM_FLAG_LOOP);
```

If a more complicated pattern of duty cycles is required that involves two sequence definitions, use [nrf\\_drv\\_pwm\\_complex\\_playback](#).

If a sequence must be updated during the playback, define a proper event handler and pass it as parameter in the call to [nrf\\_drv\\_pwm\\_init](#). By default, the handler is called at the end of the playback (or right before the next iteration in a loop), but it can be configured to be called when a sequence is completely loaded.

For example:

```

static void event_handler(nrf_drv_pwm_evt_type_t event_type)
{
    if (event_type == NRF_DRV_PWM_EVT_END_SEQ0)
    {
        // Update sequence 0.
        /* ... */
    }
    else if (event_type == NRF_DRV_PWM_EVT_END_SEQ1)
    {
        // Update sequence 1.
        /* ... */
    }
}

/* ... */

err_code = nrf_drv_pwm_init(&m_pwm0, &config0, event_handler);

/* ... */

```

```
nrf_drv_pwm_complex_playback(&m_pwm0, &seq0, &seq1, 1,
    NRF_DRV_PWM_FLAG_LOOP |
    NRF_DRV_PWM_FLAG_SIGNAL_END_SEQ0 |
    NRF_DRV_PWM_FLAG_SIGNAL_END_SEQ1 |
    NRF_DRV_PWM_FLAG_NO_EVT_FINISHED);
```

Entries in the array of duty cycle values can be modified directly, since the peripheral loads the values as the sequence is played back. If the whole sequence definition must be changed, use [nrf\\_drv\\_pwm\\_sequence\\_update](#). If only some part of the definition must be modified, you can use one of the other updating functions.

At any time, the playback can be stopped by calling [nrf\\_drv\\_pwm\\_stop](#). The pulse generation will stop at the end of the current PWM period.

When the PWM driver is no longer needed or its configuration must be changed, call [nrf\\_drv\\_pwm\\_uninit](#).

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

9. nRF5 SDK v11.0.0-2.alpha

## QDEC

The quadrature decoder (QDEC) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the QDEC. See the API documentation for the [QDEC HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [QDEC driver](#) for details.

The [QDEC Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

10. nRF5 SDK v11.0.0-2.alpha

## RNG

The random number generator (RNG) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the random number generator. See the API documentation for the [RNG HAL and driver](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [RNG driver](#) for details.

The [Random Number Generator Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

11. nRF5 SDK v11.0.0-2.alpha

## SAADC

*This information applies to the nRF52 Series only.*

The SAADC driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the SAADC peripheral. See the API documentation for the [SAADC HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [SAADC driver](#) for details.

Key features include:

- Blocking function to convert one sample on a single channel.
- Non-blocking function for triggering conversion on all enabled channels to a provided buffer.
- Double buffering: you can set up two buffers, and conversion to the new buffer starts immediately after the first one is filled.
- PAN-28 (scan mode not functional) workaround inside the driver: scan mode is emulated inside the driver, resulting in an interrupt being generated at every conversion on every channel.
- Optimal performance with interrupt only at the end of buffer conversion is achieved only on a single channel due to PAN-28.

## Driver configuration

The SAADC default configuration is located in `nrf_drv_config.h`. The SAADC must be explicitly enabled in `nrf_drv_config.h` before the driver can be used. If [nrf\\_drv\\_saadc\\_init](#) is called with a NULL pointer to the configuration structure, the default configuration from `nrf_drv_config.h` is used. You must provide an event handler function during initialization.

Driver configuration includes:

- Resolution
- Enabling or disabling oversampling
- Interrupt priority level

#### Note

If oversampling is enabled, only one channel can be enabled. [nrf\\_drv\\_saadc\\_channel\\_init](#) asserts if this condition is not met.

Each channel is configured independently. Channels are configured and enabled by the function [nrf\\_drv\\_saadc\\_channel\\_init](#). [NRF\\_DRV\\_SAADC\\_DEFAULT\\_CHANNEL\\_CONFIG\\_SE](#) or [NRF\\_DRV\\_SAADC\\_DEFAULT\\_CHANNEL\\_CONFIG\\_DIFFERENTIAL](#) can be used to build a default configuration for a channel.

Example:

```
nrf_saadc_channel_config_t channel_config =  
    NRF_DRV_SAADC_DEFAULT_CHANNEL_CONFIG_SE(NRF_SAADC_INPUT_AIN6);  
channel_config.gain = NRF_SAADC_GAIN1_2;  
//Initialization and enabling of channel 0 to use analog input 6.  
err_code = nrf_drv_saadc_channel_init(0, &channel_config);
```

[nrf\\_drv\\_saadc\\_gpio\\_to\\_ain](#) can be used to convert a GPIO pin number to an analog input number.

## Using the SAADC driver

The SAADC driver can be used in blocking mode or non-blocking mode.

### Blocking mode

The function [nrf\\_drv\\_saadc\\_sample\\_convert](#) is blocking and returns when the requested channel is sampled. The channel must be initialized before it can be used. If the SAADC is busy, the function returns with an error. In blocking mode, the driver does not use the peripheral interrupt and there is no context switching inside the driver.

### Non-blocking mode

The function [nrf\\_drv\\_saadc\\_buffer\\_convert](#) can be used to start conversion in non-blocking mode. The function returns immediately after the buffer is configured. If the driver is busy, it returns with an error. [nrf\\_drv\\_saadc\\_buffer\\_convert](#) sets the SAADC up for conversion, but does not trigger sampling. To trigger sampling, call the function [nrf\\_drv\\_saadc\\_sample](#) or, through PPI, use the task SAMPLE from SAADC. [nrf\\_drv\\_saadc\\_task\\_address\\_get](#) can be used to get the task address. Single sampling triggers conversion on all initialized channels. When the requested buffer is filled with samples, an event of type [NRF\\_DRV\\_SAADC\\_EVT\\_DONE](#) is generated.

Example for sampling triggered by the CPU:

```
err_code = nrf_drv_saadc_buffer_convert(buffer, 1);
//Check error.
err_code = nrf_drv_saadc_sample();
//Check error.
//Event handler is called immediately after conversion is finished.
```

Example for setting up sampling by PPI:

```
err_code = nrf_drv_saadc_buffer_convert(buffer, length);
//Check error.
uint32_t sample_task = nrf_drv_saadc_task_address_get
(NRF_SAADC_TASK_SAMPLE);
//Set task end point in PPI using sample_task variable.
```

The driver supports double buffering, which means that [nrf\\_drv\\_saadc\\_buffer\\_convert](#) can be called twice and the buffer that is provided in the second call will be used immediately after the first one is filled.

Continuous conversion can be achieved by setting up two buffers and calling [nrf\\_drv\\_saadc\\_buffer\\_convert](#) again in the event handler to switch between them later.

Example for setting up two buffers:

```
err_code = nrf_drv_saadc_buffer_convert(buffer1, length1);
//Check error.
err_code = nrf_drv_saadc_buffer_convert(buffer2, length2);
//Check error.
```

Example for setting up a completed buffer again in the event handler:

```
void event_handler(nrf_drv_saadc_evt_t * p_event)
{
    if (p_event->type == NRF_DRV_SAADC_EVT_DONE)
    {
        err_code = nrf_drv_saadc_buffer_convert
(p_event->data.done.p_buffer, SAMPLES_IN_BUFFER);
        //Check error.
    }
}
```

To configure the SAADC to continue conversion to the second buffer, SAADC interrupts must be handled. However, if only one channel is used, EasyDMA could automatically start processing using the second buffer if PPI is used to bind SAADC END events with the SAADC START task. This is not done by the driver because it is only controlling the SAADC peripheral, but it can be done in the application code.

## Limits

The driver can generate events of type [NRF\\_DRV\\_SAADC\\_EVT\\_LIMIT](#) if the converted sample on a given channel is exceeding a limit. You can set two limits per channel: high and low. The function [nrf\\_drv\\_saadc\\_limit\\_set](#) can be used to configure these limits. The function sets both limits; if only one limit should be enabled, the second can be disabled by using [NRF\\_DRV\\_SAADC\\_LIMITH\\_DISABLED](#) for the high limit or [NRF\\_DRV\\_SAADC\\_LIMITL\\_DISABLED](#) for the low limit.

The limit event will be generated by every sample that exceeds the limit. If this is not desired, it can be disabled in the event handler.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

12. nRF5 SDK v11.0.0-2.alpha

## SPI master

The SPI master driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the SPI and SPIM peripherals. See the API documentation for the [SPI HAL](#) and [SPIM HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [SPI master driver](#) for details.

Key features include:

- Multi-instance support.
- Common API for two types of peripherals: legacy peripherals (SPI) and peripherals using EasyDMA (SPIM). The type of peripheral that is used can be configured independently for each driver instance in `nrf_drv_config.h`.
- Two modes of data transfer: blocking and non-blocking (with end notification done via an event handler that is provided during initialization).
- EasyDMA and legacy support (SPIM and SPI).
- Supports triggering transfer from PPI (SPIM only).
- Supports post-incrementation of buffer addresses and repeated transfers (SPIM only).

Note that peripherals using EasyDMA can work only with buffers that are placed in the Data RAM region. Under certain circumstances, compilers might choose to use a different region for data placement and, for example, place a constant buffer in the code FLASH. In such a case, the SPIM peripheral cannot be used to transfer data from the buffer.

To quickly get started, use the sample code from the [SPI Master Example](#).

## Driver configuration

The SPI master driver can use multiple instances of the SPI/SPIM peripherals, and it provides a common API for both peripheral types. The instances of the peripherals that are to be assigned to the driver must be selected statically in `nrf_drv_config.h`. In the same way, you configure whether a given instance uses EasyDMA (thus whether SPIM or SPI is used).

For example:

```
#define SPI0_ENABLED 1
#define SPI0_USE_EASY_DMA 0
```

For enabled instances, you must provide the default configuration settings that are used by the [NRF\\_DRV\\_SPI\\_DEFAULT\\_CONFIG](#) macro. These settings are also defined in `nrf_drv_config.h`, for example:

```
#define SPI0_CONFIG_SCK_PIN          2
#define SPI0_CONFIG_MOSI_PIN        3
#define SPI0_CONFIG_MISO_PIN        4
#define SPI0_CONFIG_IRQ_PRIORITY    APP_IRQ_PRIORITY_LOW
```

## Basic usage

Call the [NRF\\_DRV\\_SPI\\_INSTANCE](#) macro with the ID of the peripheral instance that you want to use to create a driver instance. For example, this code will create a driver instance that uses the SPI0 or SPIM0 peripheral (depending on the value assigned to `SPI0_CONFIG_USE_EASY_DMA` in `nrf_drv_config.h`):

```
static const nrf_drv_spi_t m_spi_master_0 = NRF_DRV_SPI_INSTANCE(0);
```

Next, call [nrf\\_drv\\_spi\\_init](#) to initialize and configure the instance. See the [nrf\\_drv\\_spi\\_config\\_t](#) structure documentation for possible configuration options.

For example:

```
uint32_t err_code;

nrf_drv_spi_config_t config = NRF_DRV_SPI_DEFAULT_CONFIG(0);
config.frequency = NRF_DRV_SPI_FREQ_1M;
config.mode      = NRF_DRV_SPI_MODE_3;
config.bit_order = NRF_DRV_SPI_BIT_ORDER_LSB_FIRST;

err_code = nrf_drv_spi_init(&m_spi_master_0, &config, NULL);
if (err_code != NRF_SUCCESS)
{
    // Initialization failed. Take recovery action.
}
```

After successful initialization, you can call [nrf\\_drv\\_spi\\_transfer](#) to request SPI transfers.

The SPI master driver can automatically control the Slave Select signal (see [nrf\\_drv\\_spi\\_config\\_t::ss\\_pin](#)). This feature can only be used if there is exactly one slave. If there are multiple slaves, the Slave Select signal must be controlled externally.

When the SPI master driver instance is no longer needed or its configuration must be changed, call [nrf\\_drv\\_spi\\_uninit](#).

## Advanced usage

In non-blocking mode, you can use the [nrf\\_drv\\_spi\\_xfer](#) function that supports complex transfers. As parameters, [nrf\\_drv\\_spi\\_xfer](#) takes a pointer to a transfer descriptor ([nrf\\_drv\\_spi\\_xfer\\_desc\\_t](#)) and flags that enable more features.

The following sections show complex scenarios that use different transfer types. These scenarios are not supported with SPI, but require SPIM.

### Starting a transfer from PPI

You can use [nrf\\_drv\\_spi\\_xfer](#) to set up a transfer and start it synchronously using PPI. To do so, specify the [NRF\\_DRV\\_SPI\\_FLAG\\_HOLD\\_XFER](#) flag and use [nrf\\_drv\\_spi\\_start\\_task\\_get](#) to get the address of the task that starts the transfer.

The following example shows how to configure the driver to setup a transfer without starting it:

```
nrf_drv_spi_xfer_desc_t xfer = NRF_DRV_SPI_XFER_TRX(p_tx_data,
    tx_length, p_rx_data, rx_length);
ret_code_t ret = nrf_drv_spi_xfer(&spi, &xfer,
    NRF_DRV_SPI_FLAG_HOLD_XFER);
// SPIM is now configured and ready to be started.
if (ret == NRF_SUCCESS)
{
    uint32_t start_tsk_addr = nrf_drv_spi_start_task_get(&spi);
    // Set up PPI to trigger the transfer.
}
```

### Repeated transfers

You can configure the SPI driver to perform a certain amount of PPI-triggered transfers. To do so, specify the [NRF\\_DRV\\_SPI\\_FLAG\\_REPEATED\\_XFER](#) flag. The driver will not put the instance into busy state in this case, because in this mode, you control when the sequence of transfers is complete. Use the END event to count the number of completed transfers. To get the address of the END event, use [nrf\\_drv\\_spi\\_end\\_event\\_get](#).

Specify the [NRF\\_DRV\\_SPI\\_FLAG\\_TX\\_POSTINC](#) and [NRF\\_DRV\\_SPI\\_FLAG\\_RX\\_POSTINC](#) flags to enable post-incrementation of buffer addresses.



Additionally, you can specify the [NRF\\_DRV\\_SPI\\_FLAG\\_NO\\_XFER\\_EVT\\_HANDLER](#) flag to disable calling the user event handler after each transfer completion. If you specify this flag, transfer completion interrupts are disabled if EasyDMA and SPIM is used.

The following example shows how to set up a sequence of SPI transfers that will be triggered using PPI (for example, using RTC or TIMER compare events). A single transfer consists of a transmission of 1 byte and a reception of 3 bytes. The TX data is repeated for all transfers, but the RX data from all transfers is collected in the buffer. Therefore, [NRF\\_DRV\\_SPI\\_FLAG\\_RX\\_POSTINC](#) is set so that the buffer address for the received data is incremented. The end of sequence is controlled externally, and the processing of data is postponed until the end of the sequence (the user event handler is disabled).

```
nrf_drv_spi_xfer_desc_t xfer = NRF_DRV_SPI_XFER_TRX(p_tx_data, 1,
    p_rx_data, 3);
uint32_t flags = NRF_DRV_SPI_FLAG_HOLD_XFER          |
    NRF_DRV_SPI_FLAG_RX_POSTINC                      |
    NRF_DRV_SPI_FLAG_NO_XFER_EVT_HANDLER            |
    NRF_DRV_SPI_FLAG_REPEATED_XFER;

ret_code_t ret = nrf_drv_spi_xfer(&spi, &xfer, flags);
// SPIM is now configured and ready to be started.
if (ret == NRF_SUCCESS)
{
    uint32_t start_tsk_addr = nrf_drv_spi_start_task_get(&spi);
    // Set up PPI to trigger the transfer.
    uint32_t end_evt_addr = nrf_drv_spi_end_event_get(&spi);
    // Set up PPI to count the number of transfers.
}
```

In this example, there is no interrupt from SPI unless an error condition is detected.

## Events

When a transfer is complete, which means that the requested amount of data was transferred or an error condition was detected, the driver generates an event (unless this was suppressed with the [NRF\\_DRV\\_SPI\\_FLAG\\_NO\\_XFER\\_EVT\\_HANDLER](#) flag). This event contains the transfer descriptor.

It is possible to start another transfer from the event handler context.

There is no context blocking in the interrupt handler. Therefore, it is assumed that the driver API will not be called from a context with higher priority than the instance interrupt, because this might cause driver failure.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

13. nRF5 SDK v11.0.0-2.alpha

## SPI slave

The SPI slave driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the SPIS peripheral. See the API documentation for the [SPIS HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [SPI slave driver](#) for details.

Key features include:

- Multi-instance support
- SPI slave EasyDMA functionality for reading from and writing to the RAM

Note that peripherals using EasyDMA can work only with buffers that are placed in the data RAM region. Under certain circumstances, compilers might choose to use a different region for data placement and, for example, place a constant buffer in the code FLASH. In such a case, the SPIS peripheral cannot be used to transfer data from the buffer.

To quickly get started, use the sample code from the [SPI Slave Example](#).

## Driver configuration

The SPI slave driver can use multiple instances of SPIS. The instances of the peripherals that are to be assigned to the driver must be selected statically in `nrf_drv_config.h`, for example:

```
#define SPIS1_ENABLED 1
```

For enabled instances, you must provide the default configuration settings that are used by the [NRF\\_DRV\\_SPIS\\_DEFAULT\\_CONFIG](#) macro. These settings are also defined in `nrf_drv_config.h`, for example:

```
#define SPIS1_CONFIG_SCK_PIN      2
#define SPIS1_CONFIG_MOSI_PIN    3
#define SPIS1_CONFIG_MISO_PIN    4
#define SPIS1_CONFIG_IRQ_PRIORITY APP_IRQ_PRIORITY_LOW
```

## Using the SPI slave driver

Call the [NRF\\_DRV\\_SPIS\\_INSTANCE](#) macro with the ID of the peripheral instance that you want to use to create a driver instance. For example, this code will create a driver instance that uses the SPIS1 peripheral:

```
static const nrf_drv_spis_t m_spi_slave_1 = NRF_DRV_SPI_INSTANCE(1);
```

Next, call [nrf\\_drv\\_spis\\_init](#) to initialize and configure the instance. See the [nrf\\_drv\\_spis\\_config\\_t](#) structure documentation for possible configuration options. To receive data, the SPI buffers must be set by calling [nrf\\_drv\\_spis\\_buffers\\_set](#) after initialization.

For example:

```
static volatile uint8_t m_slave_tx_buffer[8];
static volatile uint8_t m_slave_rx_buffer[8];

static void spi_slave_handler(nrf\_drv\_spis\_event\_t event)
{
    if (event.evt_type == NRF\_DRV\_SPIS\_XFER\_DONE)
    {
        // Received bytes: event.rx_amount.
    }
}

/* ... */

uint32_t err_code;

nrf\_drv\_spis\_config\_t config = NRF\_DRV\_SPIS\_DEFAULT\_CONFIG(1);
config.mode = NRF\_DRV\_SPIS\_MODE\_3;
config.bit_order = NRF\_DRV\_SPIS\_BIT\_ORDER\_LSB\_FIRST;

err_code = nrf\_drv\_spis\_init(&m_spi_slave_1, &config,
    spi_slave_handler);
if (err_code != NRF_SUCCESS)
{
    // Initialization failed. Take recovery action.
}

err_code = nrf\_drv\_spis\_buffers\_set(mp_spis,
    (uint8_t*) m_slave_tx_buffer,
    sizeof(m_slave_tx_buffer),
    (uint8_t*) m_slave_rx_buffer,
    sizeof(m_slave_rx_buffer));
if (err_code != NRF_SUCCESS)
{
    // Buffer setup failed. Take recovery action.
}
```

The callback handler function will be called after every successful buffer setup and every completed transfer. New buffers must be set by calling [nrf\\_drv\\_spis\\_buffers\\_set](#) after every finished transaction. Otherwise, the transaction is ignored, and the default character is clocked out.

When the SPI slave driver instance is no longer needed or its configuration must be changed, call [nrf\\_drv\\_spis\\_uninit](#).

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

14. nRF5 SDK v11.0.0-2.alpha

## SWI

The SWI driver provides functions to manage software interrupts (SWI). It allows users to allocate SWIs and pass extra flags to interrupt handler functions.

### Usage

Software interrupts are allocated from the available SWI pool. By default, there are six interrupts available: SWI0 to SWI5. Note that this number is decreased if you use a SoftDevice or the Gazell or ESB protocol. For more information about resource usage for a specific protocol library, see the appropriate specification.

### Initialization

You must initialize the driver by calling the [nrf\\_drv\\_swi\\_init](#) function. If initialization is successful, the function returns [NRF\\_SUCCESS](#).

### Allocation

To allocate an interrupt, call the [nrf\\_drv\\_swi\\_alloc](#) function. As input, you must provide a pointer to the [nrf\\_swi\\_t](#) variable, an event handler function, and a value that indicates the interrupt priority. If allocation is successful, the function returns [NRF\\_SUCCESS](#), and the provided pointer variable is set.

### Event handler and flags

When allocating the driver, you must provide a callback function (see [nrf\\_swi\\_handler\\_t](#)). The arguments for this function are the [SWI](#) and some [user flags](#). The user flags are represented as bit mask.

To trigger the interrupt, call [nrf\\_drv\\_swi\\_trigger](#) with a specific flag number. When an SWI is triggered, the corresponding callback function is called with all specified flags (represented as a bit mask). The flags can be used, for example, for interrupt source tracking, or they can be ignored if not needed. Note that the callback function might be invoked with more than one flag set if the SWI is triggered more than once with different flag numbers from a same or higher priority interrupt handler.

### Special options

To disable specific SWIs from the pool, add a global define during compilation (SWI\_DISABLE0 to SWI\_DISABLE5). Any disabled SWI will not be allocated and will therefore not be available to the user.

### Example

See the following code for a usage example. Note that this example is not power optimized.

```
static nrf_swi_t my_swi;

void swi_event_handler(nrf_swi_t swi, nrf_swi_flags_t flags)
{
    /* If "my_swi" was triggered and flag #3 is present... */
    if ((swi == my_swi) && (flags & (1 << 3)))
    {
        /* Do something. */
    }
    return;
}

/*
    ....
*/

ret_code_t err_code;
nrf_swi_t my_swi;

/* Initialize the library. */
err_code = nrf_drv_swi_init();
APP_ERROR_CHECK(err_code);

/* Allocate and setup one SWI. */
err_code = nrf_drv_swi_alloc(&my_swi, swi_event_handler,
    APP_IRQ_PRIORITY_LOW);
APP_ERROR_CHECK(err_code);

while (true)
{
    /* Trigger SWI with flag #3. */
    nrf_drv_swi_trigger(my_swi, 3);

    /* Wait for a second. */
    nrf_delay_ms(1000);
}
```

## Usage with a SoftDevice

The library can be used with a SoftDevice. However, the amount of interrupts that are available to the user will be limited in this case. See the SoftDevice specification for more information about the resource usage.

Be careful when specifying APP\_IRQ\_PRIORITY\_HIGH as SWI priority. Long interrupt routines with high priority might affect the proper operation of the SoftDevice.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

15. nRF5 SDK v11.0.0-2.alpha

## Timer

The timer driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the timer. All functions in this layer are implemented as inline functions. See the API documentation for the [Timer HAL](#) for details.

The driver layer provides support for a multi-instance timer and APIs on a higher level than the HAL. See the API documentation for the [Timer driver](#) for details.

The [Timer Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

16. nRF5 SDK v11.0.0-2.alpha

## TWI master

The TWI driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the TWI registers. See the API documentation for the [TWI HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [TWI master driver](#) for details.

Key features include:

- Repeated start
- No multi-master
- Only 7-bit addressing
- Supports clock stretching
- Includes blocking and non-blocking mode
- EasyDMA and legacy support (TWIM and TWI)
- TXRX and TXTX transfers (non-blocking mode only)
- Multi-instance
- Supports triggering transfer from PPI (TWIM only)
- Supports post-incrementation of buffer addresses and repeated transfers (TWIM only)

## Initialization

Before using the instance of the driver, it must be initialized with the [nrf\\_drv\\_twi\\_init](#) function. If no event handler is provided, the instance is configured to use blocking mode. After initialization, the instance must be enabled with [nrf\\_drv\\_twi\\_enable](#).

The instance can be disabled using [nrf\\_drv\\_twi\\_disable](#) and uninitialized using [nrf\\_drv\\_twi\\_uninit](#).

## Basic usage

The TWI master driver supports two modes of operation:

- In **blocking mode**, the functions [nrf\\_drv\\_twi\\_rx](#) and [nrf\\_drv\\_twi\\_tx](#) return if the requested transfer is complete, or if an error was reported by the peripheral. In blocking mode, the driver does not use the interrupt; there is no context switch in the function.
- In **non-blocking mode**, the functions [nrf\\_drv\\_twi\\_tx](#), [nrf\\_drv\\_twi\\_rx](#), and [nrf\\_drv\\_twi\\_xfer](#) return with [NRF\\_SUCCESS](#) immediately after the transfer is set up or with [NRF\\_ERROR\\_BUSY](#) if the driver is busy (because a previous transfer is not completed). The event handler is called from the TWI interrupt context when the transfer is complete (if enabled) or if an error occurs.

The following code example for basic usage in blocking mode shows a simple write sequence:

```
uint32_t err_code;
uint8_t tx_data[] = {'a', 'b', 'c', 'd', 'e'};
const nrf_drv_twi_t twi = NRF_DRV_TWI_INSTANCE(0);

err_code = nrf_drv_twi_init(&twi, NULL, NULL);
APP_ERROR_CHECK(err_code);

nrf_drv_twi_enable(&twi);

err_code = nrf_drv_twi_tx(&twi, SLAVE_ADDRESS, tx_data, sizeof
    (tx_data), false);
APP_ERROR_CHECK(err_code);
```

## Advanced usage

In non-blocking mode, you can use the [nrf\\_drv\\_twi\\_xfer](#) function that supports complex transfers. As parameters, [nrf\\_drv\\_twi\\_xfer](#) takes the driver instance, a pointer to a transfer descriptor ([nrf\\_drv\\_twi\\_xfer\\_desc\\_t](#)), and flags that enable more features.

Note that TWI supports only the [NRF\\_DRV\\_TWI\\_FLAG\\_TX\\_NO\\_STOP](#) flag. All other flags require TWIM.

[nrf\\_drv\\_twi\\_xfer\\_desc\\_t::type](#) specifies the type of the transfer. The following transfer types are supported:

Transfer type	Description

<a href="#">NRF_DRV_TWI_XFER_TX</a>	TX transfer. By default, this transfer ends with a stop condition. Specify <a href="#">NRF_DRV_TWI_FLAG_TX_NO_STOP</a> to suppress the stop condition. Use the macro <a href="#">NRF_DRV_TWI_XFER_DESC_TX</a> to create a transfer descriptor with this type.
<a href="#">NRF_DRV_TWI_XFER_RX</a>	RX transfer with a stop condition at the end of the transfer. Use the macro <a href="#">NRF_DRV_TWI_XFER_DESC_RX</a> to create a transfer descriptor with this type.
<a href="#">NRF_DRV_TWI_XFER_TXRX</a>	TX transfer followed by an RX transfer with a repeated start condition. The transfer ends with a stop condition. Note that there is no interrupt after the TX transfer is completed (TWIM only). Use the macro <a href="#">NRF_DRV_TWI_XFER_DESC_TXRX</a> to create a transfer descriptor with this type.
<a href="#">NRF_DRV_TWI_XFER_TXTX</a>	TX transfer followed by a TX transfer with a repeated start condition. By default, the transfer ends with a stop condition. Specify <a href="#">NRF_DRV_TWI_FLAG_TX_NO_STOP</a> to suppress the stop condition. Use the macro <a href="#">NRF_DRV_TWI_XFER_DESC_TXTX</a> to create a transfer descriptor with this type.

The following sections show complex scenarios that use different transfer types. These scenarios are not supported with TWI, but require TWIM.

### Starting a transfer from PPI

You can use [nrf\\_drv\\_twi\\_xfer](#) to set up a transfer and start it synchronously using PPI. To do so, specify the [NRF\\_DRV\\_TWI\\_FLAG\\_HOLD\\_XFER](#) flag and use [nrf\\_drv\\_twi\\_start\\_task\\_get](#) to get the address of the task that starts the transfer.

The following example shows how to configure the driver to setup a TX transfer without starting it:

```
nrf_drv_twi_xfer_desc_t xfer = NRF_DRV_TWI_XFER_DESC_TX(addr, p_data,
    length);
ret_code_t ret = nrf_drv_twi_xfer(&twi, &xfer,
    NRF_DRV_TWI_FLAG_HOLD_XFER);
// TWIM is now configured and ready to be started.
if (ret == NRF_SUCCESS)
{
    uint32_t start_tsk_addr = nrf_drv_twi_start_task_get(&twi,
        xfer.type);
    // Set up PPI to trigger the transfer.
}
```

### Repeated transfers

You can configure the TWI driver to perform a certain amount of PPI-triggered transfers. To do so, specify the [NRF\\_DRV\\_TWI\\_FLAG\\_REPEATED\\_XFER](#) flag. The driver will not put the instance into busy state in this case, because in this mode, you control when the sequence of transfers is



complete. Use the STOPPED event to count the number of completed transfers. To get the address of the STOPPED event, use [nrf\\_drv\\_twi\\_stopped\\_event\\_get](#).

Specify the [NRF\\_DRV\\_TWI\\_FLAG\\_TX\\_POSTINC](#) and [NRF\\_DRV\\_TWI\\_FLAG\\_RX\\_POSTINC](#) flags to enable post-incrementation of buffer addresses.

Additionally, you can specify the [NRF\\_DRV\\_TWI\\_FLAG\\_NO\\_XFER\\_EVT\\_HANDLER](#) flag to disable calling the user event handler after each transfer completion. If you specify this flag, transfer completion interrupts are disabled if EasyDMA and TWIM is used, and the user event handler will be called only if an ERROR interrupt is triggered.

The following example shows how to set up a sequence of [NRF\\_DRV\\_TWI\\_XFER\\_TXRX](#) transfers that will be triggered using PPI (for example, using RTC or TIMER compare events). A single transfer consists of a transmission of 1 byte and a reception of 3 bytes. The TX data is repeated for all transfers, but the RX data from all transfers is collected in the buffer. Therefore, [NRF\\_DRV\\_TWI\\_FLAG\\_RX\\_POSTINC](#) is set so that the buffer address for the received data is incremented. The end of sequence is controlled externally, and the processing of data is postponed until the end of the sequence (the user event handler is disabled).

```
nrf_drv_twi_xfer_desc_t xfer = NRF_DRV_TWI_XFER_DESC_TXRX(addr,
    p_tx_data, 1, p_rx_data, 3);
uint32_t flags = NRF_DRV_TWI_FLAG_HOLD_XFER          |
    NRF_DRV_TWI_FLAG_RX_POSTINC                      |
    NRF_DRV_TWI_FLAG_NO_XFER_EVT_HANDLER            |
    NRF_DRV_TWI_FLAG_REPEATED_XFER;

ret_code_t ret = nrf_drv_twi_xfer(&twi, &xfer, flags);
// TWIM is now configured and ready to be started.
if (ret == NRF_SUCCESS)
{
    uint32_t start_tsk_addr = nrf_drv_twi_start_task_get(&twi,
        xfer.type);
    // Set up PPI to trigger the transfer.
    uint32_t stopped_evt_addr = nrf_drv_twi_stopped_event_get(&twi);
    // Set up PPI to count the number of transfers.
}
```

In this example, there is no interrupt from TWI unless an error condition is detected.

## Events

When a transfer is complete, which means that the requested amount of data was transferred or an error condition was detected, the driver generates an event (unless this was suppressed with the [NRF\\_DRV\\_TWI\\_FLAG\\_NO\\_XFER\\_EVT\\_HANDLER](#) flag). This event contains the transfer type and the transfer descriptor. If an error occurred, the length field is set to the amount of bytes that were transferred before the error occurred.

It is possible to start another transfer from the event handler context.

The functions [nrf\\_drv\\_twi\\_rx](#), [nrf\\_drv\\_twi\\_tx](#), and [nrf\\_drv\\_twi\\_xfer](#) internally disable all TWI/TWIM interrupts when setting up transfers, and then re-enable them before leaving the function. There is no context blocking in the interrupt handler. Therefore, it is assumed that the driver API will not be called from a context with higher priority than the instance interrupt, because this might cause driver failure.

The following code shows an event handler that uses TX RX with a repeated start:

```
void twi_event_handler(nrf\_drv\_twi\_evt\_t * p_event, void * p_context)
{
    if ((p_event->type == NRF\_DRV\_TWI\_EVT\_DONE) &&
        (p_event->xfer_desc.type == NRF\_DRV\_TWI\_XFER\_TXRX))
    { xfer_completed = true; }
}
```

The following code shows the transfer starting:

```
nrf\_drv\_twi\_xfer\_desc\_t xfer = NRF\_DRV\_TWI\_XFER\_DESC\_TXRX(addr,
    p_tx_buffer, tx_length, p_rx_buffer, rx_length);
ret\_code\_t ret = nrf\_drv\_twi\_xfer(&twi, &xfer, 0);
if (ret == NRF\_SUCCESS)
{
    while(xfer_completed == false){}
}
```

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

17. nRF5 SDK v11.0.0-2.alpha

## TWI slave

*This information applies to the nRF52 Series only.*

The two-wire interface slave with EasyDMA (TWIS) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the TWI. See the API documentation for the [TWIS HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [TWI slave with EasyDMA driver](#) for details.

Key features include:

- Simple events communication:

- Data request
- Transfer pending
- Transfer finished
- Error
- Two 7-bit addresses for listening
- Supports clock stretching

## Driver configuration

The following configuration options are available in `nrf_drv_config.h`:

- [`TWIS\_ASSUME\_INIT\_AFTER\_RESET\_ONLY`](#): Set this option for better performance if you activate TWI only once.
- [`TWIS\_NO\_SYNC\_MODE`](#): Set this option to always use [Asynchronous mode](#).

## Modes of operation

The TWIS driver supports two different modes: synchronous mode and asynchronous mode.

### Synchronous mode

In synchronous mode, the driver works completely without interrupts. To configure the driver to use synchronous mode, pass a NULL pointer as the pointer to the event handler during driver initialization.

In synchronous mode, events are pooled every time a state function is called. Pooling and state machine processing can be entered only once at the same time. Therefore, if a function is called from the main thread and, while it is processed, an interrupt occurs that tries to process a TWIS event again, the checking function in the interrupt is skipped.

The TWIS driver is a slave driver. Therefore, the reasons for configuring this driver to use synchronous mode are different than for a master driver. For a slave driver, it does not make sense to support sending data from the main thread and from interrupts at the same time. State testing functions like [`nrf\_drv\_twis\_is\_busy`](#) are not guaranteed to process the current state in the interrupts if the same functions are processed in the main thread. Interrupts should not wait for a state change. Otherwise, they might be locked.

The following example code shows how to use synchronous mode:

```
nrf_drv_twis_t instance = NRF_DRV_TWIS_INSTANCE(0);

int main(void)
{
    // Initialization
    bool waitingRx = false;
    nrf_drv_twis_init(&instance, NULL, NULL);
    nrf_drv_twis_enable(&instance);

    // Main loop
    while(1)
```

```

{
    // It is important to process it first
    if(waitingRx && !nrf_drv_twis_is_pending_rx(&instance))
    {
        waitingRx = false;
        // Check of errors
        uint32_t err = nrf_drv_twis_error_get(&instance);
        if(err)
        {
            // ... process errors ....
        }
        else
        {
            // ... process received data ...
        }
    }
    if(nrf_drv_twis_is_waiting_tx_buff(&instance))
    {
        // Prepare TX buffer (response for READ command)
        nrf_drv_twis_tx_prepare(&instance, txbuffer, sizeof
(txbuffer));
    }
    if(nrf_drv_twis_is_waiting_rx_buff(&instance))
    {
        // Prepare RX buffer (response for WRITE command)
        nrf_drv_twis_rx_prepare(&instance, rxbuffer, sizeof
(rxbuffer));
        waitingRx = true;
    }

    // ... The rest of the main lop tasks ...
}
}

```

## Asynchronous mode

In asynchronous mode, all events are processed in internal interrupt service runtime. When anything requires processing, the event handler is called. See [nrf\\_drv\\_twis\\_evt\\_type\\_t](#) for a list of defined events.

The event handler is always called in the following pattern: REQ -> DONE or ERROR. This means that if we get a TWIS\_EVT\_READ\_REQ message, the answer will always be TWIS\_EVT\_READ\_DONE or TWIS\_EVT\_READ\_ERROR. The answer will be sent even if there is no STOP condition on the bus, but reading finishes for example by a repeated START condition.

There are different ways to prepare buffers for receiving or sending data:

- Prepare the buffer when required inside the event handler.

- Prepare the buffer outside of the event handler by using a flag set in the event handler or by checking if the buffer is required by [nrf\\_drv\\_twis\\_is\\_waiting\\_rx\\_buff](#) or [nrf\\_drv\\_twis\\_is\\_waiting\\_tx\\_buff](#).
- Prepare the buffer before it is requested. If the buffer is not required, the [nrf\\_drv\\_twis\\_evt\\_t::buf\\_req](#) field is cleared, and the buffer can be set for the next transaction in that event.

The following pseudocode shows how to implement an event handler:

```
void twis_event_handler(nrf\_drv\_twis\_evt\_t const * const p_event)
{
    switch(p_event->type)
    {
        case TWIS\_EVT\_READ\_REQ:
            if(p_event->data.buf\_req)
            {
                // Function size_t get_buffer_to_send(char const * * const
                rptr)
                // needs to be implemented
                const char * rptr;
                size_t size = get_buffer_to_send(&rptr);
                nrf\_drv\_twis\_tx\_prepare(&m_slave_inst, rptr, size);
            }
            break;
        case TWIS\_EVT\_READ\_DONE:
            break;
        case TWIS\_EVT\_WRITE\_REQ:
            if(p_event->data.buf\_req)
            {
                // Function size_t get_buffer_to_write(char * * const wptr)
                // needs to be implemented
                char *wptr;
                size_t size = get_buffer_to_write(&wptr);
                nrf\_drv\_twis\_rx\_prepare(&m_slave_inst, wptr, size);
            }
            break;
        case TWIS\_EVT\_WRITE\_DONE:
            // Function void mark_buffer_ready(size_t amount)
            // needs to be implemented
            mark_buffer_ready(p_event->data.rx\_amount);
            break;
        default:
            break;
    }
}
```

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

18. nRF5 SDK v11.0.0-2.alpha

## UART

The UART driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the UART and UARTE peripherals. See the API documentation for the [UART HAL](#) and [UARTE HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [UART driver](#) for details.

Key features include:

- Common API for two types of peripherals: legacy peripherals (UART) and peripherals using EasyDMA (UARTE). The type of peripheral that is used can be configured either during build time or at runtime during initialization.
- Two modes of data transfer: blocking and non-blocking (with end notification done via an event handler that is provided during initialization).
- Driver supports aborting ongoing RX or TX transfers.
- Supported modes: with and without hardware flow control and with and without parity bit.
- Double RX buffering allowing to receive data continuously.

Note that peripherals using EasyDMA can work only with buffers that are placed in the Data RAM region. Under certain circumstances, compilers might choose to use a different region for data placement and, for example, place a constant buffer in the code FLASH. In such a case, the UARTE peripheral cannot be used to transfer data from the buffer.

### Driver configuration

The UART default configuration is located in `nrf_drv_config.h`. If UARTE is present on the chip, the driver can be configured at runtime to support UART mode, UARTE mode, or both. The following example shows how to configure the driver to support both modes in runtime and have legacy mode as the default:

```
#ifndef NRF52
#define UART0_CONFIG_USE_EASY_DMA false
#define UART_EASY_DMA_SUPPORT      1
#define UART_LEGACY_SUPPORT        1
#endif //NRF52
```

If only one mode is used in the application, disable the second mode in `nrf_drv_config.h` to achieve a lower memory footprint and better performance.

Call [nrf\\_drv\\_uart\\_init](#) with the `p_config` argument set to `NULL` to use the default configuration. To use a custom configuration, provide a user configuration structure, for example:

```
nrf_drv_uart_config_t config = NRF_DRV_UART_DEFAULT_CONFIG;
config.use_easy_dma = true;
ret_code = nrf_drv_uart_init(&config, NULL);
```

## Using the UART driver

### Blocking mode

If no event handler is provided during the initialization of the driver, the driver will operate in blocking mode. In this case, [nrf\\_drv\\_uart\\_tx](#) and [nrf\\_drv\\_uart\\_rx](#) will not return until the requested transfer is completed, [nrf\\_drv\\_uart\\_rx\\_abort](#) or [nrf\\_drv\\_uart\\_tx\\_abort](#) is called from a different context, or an error is detected. When an abort function is called, [nrf\\_drv\\_uart\\_tx](#) or [nrf\\_drv\\_uart\\_rx](#) returns with an error code. When an error is reported by the peripheral, [nrf\\_drv\\_uart\\_rx](#) returns with an error code.

In blocking mode, the driver does not use a peripheral interrupt, and there is no context switching inside the driver.

### Non-blocking mode

If an event handler ([nrf\\_uart\\_event\\_handler\\_t](#)) is provided during the initialization of the driver, the driver will operate in non-blocking mode. In this case, [nrf\\_drv\\_uart\\_tx](#) and [nrf\\_drv\\_uart\\_rx](#) will return immediately after transfer is started. Completion will be notified to the user by the event handler that was provided during initialization. The event handler is called in the context of the UART/UARTE interrupt. If an error is reported by the peripheral, it is also reported by the event handler. The error event contains information about the amount of data that was transferred before the error occurred.

If [nrf\\_drv\\_uart\\_rx\\_abort](#) or [nrf\\_drv\\_uart\\_tx\\_abort](#) is called, the event handler notifies that the transfer is completed, but the length field is set to the amount of data that was transferred before the abort function was called.

### Enabling RX without providing data buffer

You can enable RX without providing a buffer for incoming data. In this case, incoming data (up to 6 bytes) stays in the hardware FIFO. If more than 6 bytes are received, an overrun error condition occurs. This condition will be reported by the event handler or at the next [nrf\\_drv\\_uart\\_rx](#) call.

#### Note

This functionality is provided only for UART mode.

```
ret_code = nrf_drv_uart_init(NULL, NULL);
nrf_drv_uart_rx_enable();
//UART RX is enabled
nrf_drv_uart_rx(buffer, 2);
```

```
//UART RX is still enabled.  
nrf\_drv\_uart\_rx\_disable\(\);
```

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).

19. *nRF5 SDK v11.0.0-2.alpha*

## WDT

The watchdog timer (WDT) driver includes two layers: the hardware access layer (HAL) and the driver layer (DRV).

The hardware access layer provides basic APIs for accessing the registers of the WDT. All functions in this layer are implemented as inline functions. See the API documentation for the [WDT HAL](#) for details.

The driver layer provides APIs on a higher level than the HAL. See the API documentation for the [WDT driver](#) for details.

The [WDT Example](#) provides sample code that you can use to quickly get started.

This document was last updated on Fri Dec 18 2015.

Please send us your [feedback](#) about the documentation! For technical questions, visit the [Nordic Developer Zone](#).